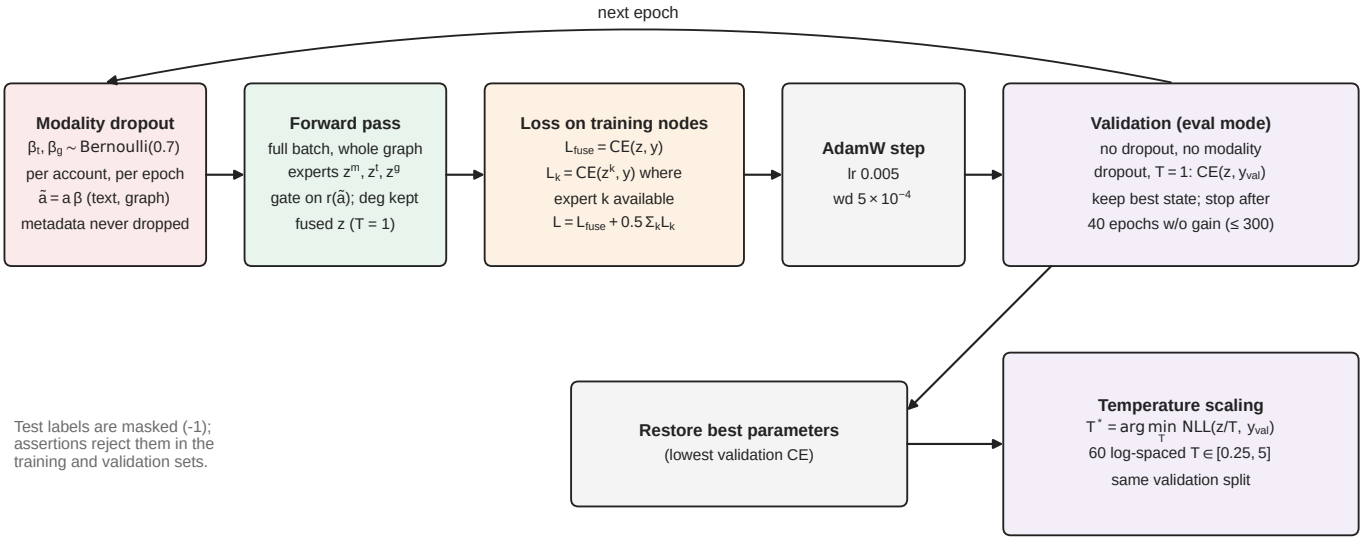
next epoch
Modality dropout
$\beta_t, \beta_g \sim$ Bernoulli(0.7)
per account, per epoch
$\tilde{a} = a\,\beta$ (text, graph)
metadata never dropped
Forward pass
full batch, whole graph
experts z^m, z^t, z^g
gate on r($\tilde{a}$); deg kept
fused z (T = 1)
Loss on training nodes
$L_{fuse} = CE(z, y)$
$L_k = CE(z^k, y)$ where
expert k available
$L = L_{fuse} + 0.5\,\Sigma_k L_k$
AdamW step
lr 0.005
wd 5×10^{-4}
Validation (eval mode)
no dropout, no modality
dropout, T = 1: CE(z, y_{val})
keep best state; stop after
40 epochs w/o gain (≤ 300)
Restore best parameters
(lowest validation CE)
Temperature scaling
$T^* = \arg\min_T$ NLL(z/T, y_{val})
60 log-spaced T ∈ [0.25, 5]
same validation split
Test labels are masked (-1);
assertions reject them in the
training and validation sets.